

## Ejercicio 2 Parte 2/2 (Código en examen)

En total son 753,664 hilos

Y la máxima capacidad de hilos por bloque es 1024

Por lo que tendríamos 736 bloques en total.

Agregaríamos al código:

num\_SM \* max\_blocks\_per\_SM y max\_threads\_per\_block

```
#include "cuda_runtime.h"
```

```
#include "device_launch_parameters.h"
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define GPUErrorAssertion(ans) { gpuAssert((ans), __FILE__, __LINE__); }
```

```
inline void gpuAssert(cudaError_t code, const char* file, int line, bool abort = true) {
```

```
    if (code != cudaSuccess) {
```

```
        fprintf(stderr, "GPUassert: %s %s %d\n", cudaGetErrorString(code), file, line);
```

```
        if (abort) exit(code);
```

```
    }
```

```
}
```

```
__global__ void imageThreshold(unsigned char* a, unsigned char* b, unsigned char*  
c, int N) {
```

```
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
```

```
    if (idx < N) {
```

```
        c[idx] = a[idx] | b[idx];
```

```
    }
```

```
}
```

```
int main() {
```

```
    const int N = 1024;
```

```
    const int dataSize = N * sizeof(unsigned char);
```

```
    unsigned char* A_cpu;
```

```
    unsigned char* B_cpu;
```

```
    unsigned char* C_cpu;
```

```
    unsigned char* A_gpu;
```

```
    unsigned char* B_gpu;
```

```
    unsigned char* C_gpu;
```

```
    A_cpu = (unsigned char*)malloc(dataSize);
```

```

B_cpu = (unsigned char*)malloc(dataSize);
C_cpu = (unsigned char*)malloc(dataSize);

for (int i = 0; i < N; i++) {
    A_cpu[i] = rand() % 127;
    B_cpu[i] = rand() % 127;
}

GPUErrorAssertion(cudaMalloc((void**)&A_gpu, dataSize));
GPUErrorAssertion(cudaMalloc((void**)&B_gpu, dataSize));
GPUErrorAssertion(cudaMalloc((void**)&C_gpu, dataSize));

GPUErrorAssertion(cudaMemcpy(A_gpu, A_cpu, dataSize,
    cudaMemcpyHostToDevice));
GPUErrorAssertion(cudaMemcpy(B_gpu, B_cpu, dataSize,
    cudaMemcpyHostToDevice));

int maxThreadsPerBlock;
cudaDeviceProp prop;
cudaGetDeviceProperties(&prop, 0);
maxThreadsPerBlock = prop.maxThreadsPerBlock;

int threadsPerBlock = maxThreadsPerBlock;
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;

imageThreshold<<<blocksPerGrid, threadsPerBlock>>>(A_gpu, B_gpu, C_gpu,
N);
GPUErrorAssertion(cudaDeviceSynchronize());

GPUErrorAssertion(cudaMemcpy(C_cpu, C_gpu, dataSize,
    cudaMemcpyDeviceToHost));

for (int i = 0; i < 10; i++) {
    printf("Result[%d] = %d\n", i, C_cpu[i]);
}

cudaFree(A_gpu);
cudaFree(B_gpu);
cudaFree(C_gpu);

free(A_cpu);
free(B_cpu);
free(C_cpu);

```

```
    return 0;  
}
```